

## Machine learning project information (30%)

### Project Details

#### **Objective:**

Each project is focused on solving a specific real-world problem through machine learning, including tasks like data preprocessing, model training, evaluation, deployment, and tracking. Students will explore the complete model-building lifecycle, from initial data analysis to deployment.

#### **Dataset:**

Each project is associated with a suitable dataset:

- **Cybersecurity Threat Detection:** CICIDS 2017 Dataset (Kaggle).
  - **Sentiment Analysis:** IMDb Movie Reviews Dataset.
  - **Plant Disease Detection:** Plant Village Dataset.
  - **Patient Health Risk Prediction:** Diabetes Health Indicators Dataset or Heart Disease UCI Dataset.
- 

### Project Components

#### 1. Feature Engineering and Exploratory Data Analysis (EDA)

- **Objective:** Understand the dataset, identify trends, and engineer new features to improve model performance.
- **Tasks:**
  - Conduct an exploratory data analysis to observe feature distributions, correlations, and patterns.
  - Visualize data with histograms, box plots, and heatmaps.
  - Create additional features if they could enhance model accuracy.
- **Deliverable:** Insights from EDA, supported by visualizations, included in the final report.

#### 2. Model Training and Comparison

- **Objective:** Train, compare, and select the best-performing model, exploring ensemble methods if applicable.
- **Tasks:**

## MOD10: Machine learning

Fall 2025

Instructor: Mohammed A. Shehab

- Train at least two different models (e.g., decision tree, random forest) and evaluate their performance.
- Experiment with ensemble techniques (e.g., voting, stacking) to potentially improve performance.
- **Deliverable:** A report section discussing model comparison, choice, and the effect of any ensemble methods.

### 3. Evaluation Metrics

- **Objective:** Assess model performance through a variety of metrics to understand strengths and weaknesses.
- **Tasks:**
  - For classification: Calculate metrics like ROC, AUC, accuracy, precision, recall, and F1-score.
  - For regression (if relevant): Use R-squared, MAE, and MSE.
- **Deliverable:** Comparative analysis of metrics, documented in the report.

### 4. Interpretability and Explainability

- **Objective:** Analyze the model's decision-making process to understand feature importance and enhance transparency.
- **Tasks:**
  - Use interpretability tools (e.g., SHAP, LIME) to explain key features influencing predictions, especially for complex models.
- **Deliverable:** A report section explaining the most influential features based on model interpretations.

### 5. Deployment with Docker

- **Objective:** Containerize the model using Docker and deploy it with a REST API for real-time predictions.
- **Tasks:**
  - Create a Dockerfile with all dependencies, enabling easy deployment.
  - Set up a REST API (using Flask or FastAPI) to serve predictions.
- **Deliverable:** Docker image with instructions for running and testing the container.

### 6. Performance Tracking with MLflow

- **Objective:** Log model metrics, parameters, and experiment versions using MLflow for performance tracking.

**MOD10: Machine learning****Fall 2025****Instructor: Mohammed A. Shehab**○ **Tasks:**

- Track model metrics during training and deployment for analysis and comparison.
  - Use MLflow to log all experiments, parameters, and versions.
  - **Deliverable:** MLflow logs showing metrics and comparison across model versions.
- 

## Project Workflow

**1. Data Preprocessing and EDA**

- Load and clean the dataset, handling missing values, normalizing or scaling numerical features, and encoding categorical variables.
- Conduct EDA, visualizing feature distributions and identifying potential relationships to guide feature engineering.

**2. Model Training, Comparison, and Ensemble**

- Train multiple models (e.g., logistic regression, decision trees) and evaluate them using relevant metrics.
- Experiment with ensemble methods (e.g., voting, stacking) if it improves model performance.

**3. Evaluation Metrics and Analysis**

- Evaluate models on various metrics, comparing results to understand model strengths and weaknesses, particularly for imbalanced datasets.

**4. Interpretability and Explainability**

- Use tools like SHAP or LIME to explain model decisions and analyze feature importance, ensuring model transparency.

**5. Docker Deployment**

- Build and deploy the model in a Docker container with a REST API, allowing for easy real-time predictions.

**6. MLflow Tracking**

- Log all experiments, tracking parameters and metrics over multiple model versions with MLflow.
-

## Deliverables

- 1. Code:**
    - Scripts or Jupyter notebooks for EDA, model training, evaluation, deployment, and MLflow tracking.
  - 2. EDA Visualizations:**
    - Visualizations such as histograms, box plots, and heatmaps documenting data characteristics and feature correlations.
  - 3. Model Comparison and Interpretability Analysis:**
    - Report sections analyzing model performance and feature importance, including any ensemble techniques applied.
  - 4. Docker and MLflow:**
    - Dockerfile and container instructions for deployment, plus MLflow logs tracking model metrics.
  - 5. Final Report:**
    - A report summarizing project objectives, methodology, EDA findings, model evaluation, interpretability analysis, deployment steps, and key takeaways.
- 

## Evaluation Criteria

- 1. Data Preprocessing and EDA (20%)**
  - **EDA Insights (10%):** Effective visualizations and insights on feature distributions and correlations.
  - **Feature Engineering (10%):** Creation of new features and analysis of their impact on model performance.
- 2. Model Training, Comparison, and Ensemble (25%)**
  - **Model Implementations (15%):** Successful training and tuning of at least two models, with clear model selection rationale.
  - **Ensemble Techniques (10%):** Exploration and analysis of ensemble techniques, if applied.
- 3. Evaluation Metrics (15%)**
  - **Metric Analysis (15%):** Calculation and interpretation of metrics, with discussion on metric effectiveness for the project's objectives.

**MOD10: Machine learning**

**Fall 2025**

**Instructor: Mohammed A. Shehab**

**4. Interpretability and Explainability (10%)**

- **Feature Importance (10%):** Clear analysis using SHAP, LIME, or similar tools to explain key features affecting model predictions.

**5. Deployment and Performance Tracking (20%)**

- **Docker Deployment (10%):** Complete Docker container with functional API for predictions.
- **MLflow Tracking (10%):** Accurate MLflow logging of parameters, metrics, and model versions.

**6. Presentation (10%)**

- **Clarity and Organization (5%):** Clear presentation, effectively communicating objectives, methodology, and results.
- **Visuals and Engagement (5%):** Use of visuals for data insights and effective collaboration among group members.